










small6 (36.5B)

large6 (203B)

all, 21 judges (367.5B+, 3 paid APIs)

Chatbot Arena,  $n_L = 2000$   
excess held-out log lossMT-Bench,  $n_L = 160$   
excess held-out log lossPandaLM,  $n_L = 100$   
excess held-out log losshuman calibration labels  $n_H$ human calibration labels  $n_H$ human calibration labels  $n_H$ 

Cons-Cal

DIAL- $\mu$