










small6 (36.5B)

large6 (203B)

all, 21 judges (367.5B+, 3 paid APIs)

Chatbot Arena, $n_H = 300$
excess held-out log lossMT-Bench, $n_H = 80$
excess held-out log lossPandaLM, $n_H = 60$
excess held-out log lossLLM comparisons n_L LLM comparisons n_L LLM comparisons n_L

--- Human -▼- Cons-Cal -●- DIAL- μ